
Constraint Satisfaction Problems

AIMA: Chapter 6


Constraint Satisfaction Problems

A CSP consists of:

- *Finite set of variables* X_1, X_2, \dots, X_n
- *Nonempty domain of possible values* for each variable D_1, D_2, \dots, D_n *where* $D_i = \{v_1, \dots, v_k\}$
- *Finite set of constraints* C_1, C_2, \dots, C_m
 - Each *constraint* C_i limits the values that variables can take, e.g., $X_1 \neq X_2$ A *state* is defined as an *assignment* of values to some or all variables.
- A *consistent assignment* does not violate the constraints.
- **Example: Sudoku**

Example: Cryptarithmic

$$\begin{array}{r}
 X_3 \quad X_2 \quad X_1 \\
 \quad T \quad W \quad O \\
 + \quad T \quad W \quad O \\
 \hline
 F \quad O \quad U \quad R
 \end{array}$$



- **Variables:** $F T U W R O, X_1 X_2 X_3$
- **Domain:** $\{0,1,2,3,4,5,6,7,8,9\}$
- **Constraints:**
 - $Alldiff(F, T, U, W, R, O)$
 - $O + O = R + 10 \cdot X_1$
 - $X_1 + W + W = U + 10 \cdot X_2$
 - $X_2 + T + T = O + 10 \cdot X_3$
 - $X_3 = F, T \neq 0, F \neq 0$

Constraint satisfaction problems


- An assignment is *complete* when every variable is assigned a value.
- A *solution* to a CSP is a complete assignment that satisfies all constraints.
- Applications:
 - Map coloring
 - Line Drawing Interpretation
 - Scheduling problems
 - Job shop scheduling
 - Scheduling the Hubble Space Telescope
 - Floor planning for VLSI
- Beyond our scope: CSPs that require a solution that maximizes an *objective function*.

Example: Map-coloring



- **Variables:** *WA, NT, Q, NSW, V, SA, T*
- **Domains:** $D_i = \{\text{red, green, blue}\}$
- **Constraints:** adjacent regions must have different colors
 - e.g., $WA \neq NT$
 - So (WA, NT) must be in $\{(\text{red, green}), (\text{red, blue}), (\text{green, red}), \dots\}$

Example: Map-coloring




Solutions are **complete** and **consistent** assignments,

- e.g., WA = red, NT = green, Q = red, NSW = green,
V = red, SA = blue, T = green

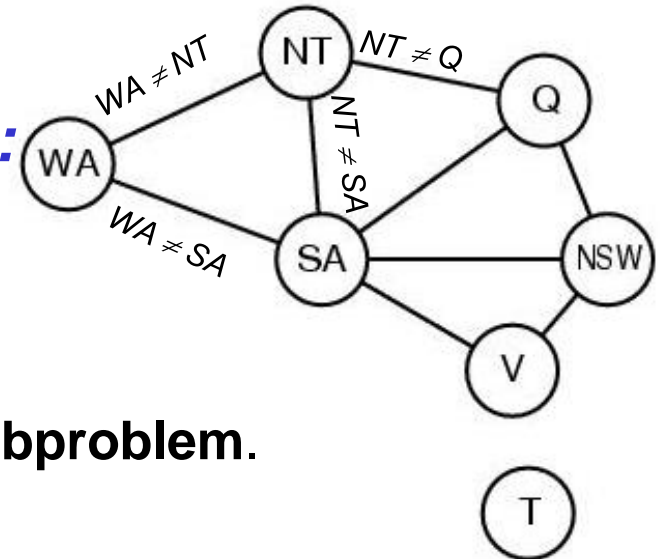
Benefits of CSP

- **Clean specification of many problems, generic goal, successor function & heuristics**
 - Just represent problem as a CSP & solve with general package
- **CSP “knows” which variables violate a constraint**
 - And hence where to focus the search
- **CSPs: Automatically prune off all branches that violate constraints**
 - (State space search could do this only by *hand-building constraints into the successor function*)

CSP Representations



- **Constraint graph:**
 - *nodes* are variables
 - *arcs* are constraints
- **Standard representation pattern:**
 - variables with values
- **Constraint graph** simplifies search.
 - e.g. Tasmania is an independent subproblem.
- **This problem: A binary CSP:**
 - each constraint relates two variables



Varieties of CSPs

- **Discrete variables**

- finite domains:
 - n variables, domain size $d \rightarrow O(d^n)$ complete assignments
 - e.g., Boolean CSPs, includes Boolean satisfiability (NP-complete)
 - **Line Drawing Interpretation**
- infinite domains:
 - integers, strings, etc.
 - e.g., job scheduling, variables are start/end days for each job
 - need a constraint language, e.g., $StartJob_1 + 5 \leq StartJob_3$

- **Continuous variables**

- e.g., start/end times for Hubble Space Telescope observations
- linear constraints solvable in polynomial time by linear programming

Varieties of constraints

- **Unary** constraints involve a single variable,
 - e.g., $SA \neq \text{green}$
- **Binary** constraints involve pairs of variables,
 - e.g., $SA \neq WA$
- **Higher-order** constraints involve 3 or more variables
 - e.g., crypt-arithmetic column constraints
- **Preference** (soft constraints) e.g. *red is better than green* can be represented by a cost for each variable assignment
 - Constrained optimization problems.


Idea 1: CSP as a search problem

- **A CSP can easily be expressed as a search problem**
 - *Initial State*: the empty assignment {}.
 - *Successor function*: Assign value to any unassigned variable *provided that there is not a constraint conflict*.
 - *Goal test*: the current assignment is complete.
 - *Path cost*: a constant cost for every step.
- **Solution is always found at depth n , for n variables**
 - Hence Depth First Search can be used

Backtracking search



- Note that variable assignments are *commutative*
 - Eg [step 1: **WA = red**; step 2: **NT = green**]
equivalent to [step 1: **NT = green**; step 2: **WA = red**]
 - Therefore, a *tree search*, not a *graph search*
- Only need to consider assignments to a single variable at each node
 - $b = d$ and there are d^n leaves
- Depth-first search for CSPs with single-variable assignments is called *backtracking* search
- Backtracking search is the basic *uninformed* algorithm for CSPs
- Can solve *n*-queens for $n \approx 25$

Backtracking example



And so on....

Backtracking example




And so on....

Idea 2: Improving backtracking efficiency


- **General-purpose** methods & heuristics can give huge gains in speed, **on average**
- **Heuristics:**
 - Q: Which variable should be assigned next?
 1. **Most constrained** variable
 2. **Most constraining** variable
 - Q: In what order should that variable's values be tried?
 3. **Least constraining** value
 - Q: Can we detect inevitable failure early?
 4. **Forward checking**

Heuristic 1: Most constrained variable

- Choose a variable with the *fewest legal values*





- a.k.a. *minimum remaining values (MRV)* heuristic



Heuristic 2: Most constraining variable

- Tie-breaker among most constrained variables
- Choose the variable with the *most constraints on remaining variables*



Heuristic 3: Least constraining value


- Given a variable, *choose the least constraining value*:
 - the one that rules out the fewest values in the remaining variables



Note: demonstrated here independent of the other heuristics



Heuristic 4: Forward checking

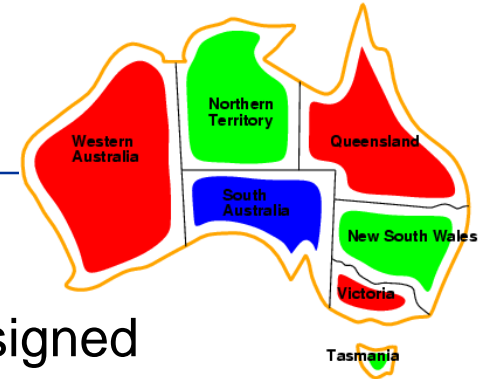


- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values, *given its neighbors*




- (For later: Edge & Arc consistency are variants)


Forward checking



- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



Forward checking




- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values

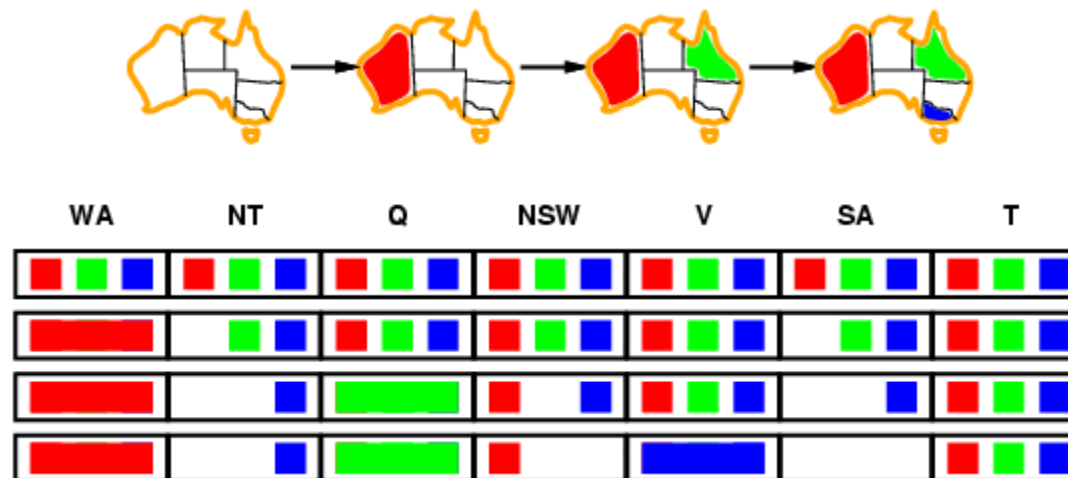


WA	NT	Q	NSW	V	SA	T
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

Forward checking




- **Idea:**
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



- **A Step toward AC-3: The most efficient algorithm**



Example: 4-Queens Problem

	X1	X2	X3	X4
1				
2				
3				
4				





(From Bonnie Dorr, U of Md, CMSC 421)



Example: 4-Queens Problem





Example: 4-Queens Problem



Example: 4-Queens Problem




Example: 4-Queens Problem





Example: 4-Queens Problem

Picking up a little later after two steps of backtracking....



	1	2	3	4
1		●		
2	★	●	●	●
3		●		
4			●	





Example: 4-Queens Problem





Example: 4-Queens Problem





Example: 4-Queens Problem





Example: 4-Queens Problem




Example: 4-Queens Problem



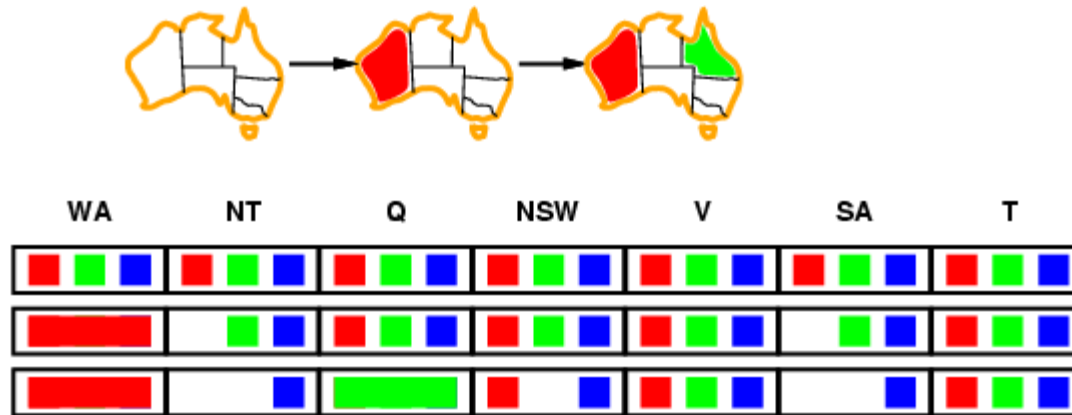
Example: 4-Queens Problem



Towards Constraint propagation




- Forward checking propagates information from *assigned* to *unassigned* variables, but doesn't provide early detection for all failures:




- NT and SA cannot both be blue!
- **Constraint propagation** goes beyond forward checking & repeatedly enforces constraints locally

Interpreting line drawings and the invention of constraint *propagation* algorithms



We Interpret Line Drawings As 3D


- We have strong intuitions about line drawings of simple geometric figures:
- We naturally interpret 2D line drawings as planar *representations* of *3D* objects.
- We interpret each line as being either a *convex*, *concave* or *occluding edge* in the actual object.



Interpretation as Convexity Labeling

Each edge in an image can be interpreted to be either *a convex edge*, *a concave edge* or *an occluding edge*:


- **+** labels a *convex edge* (angled toward the viewer);
- **-** labels a *concave edge* (angled away from the viewer);
- **→** labels an *occluding edge*. To its right is the body for which the arrow line provides an edge. On its left is space.



Huffman/Clowes


Line Drawing Interpretation

- ***Given:*** a line drawing of a simple “blocks world” physical image
- ***Compute:*** a set of junction labels that yields a consistent physical interpretation




Huffman/Clowes Junction Labels


- A simple trihedral image can be automatically interpreted given only information about each *junction* in the image.
- Each interpretation gives *convexity* information for each junction.
- This interpretation is based on the *junction type*. (All junctions involve at most *three* lines.)




arrow
junction



Y
junction





L
junction



T
junction


Arrow Junctions have *only* 3 interpretations



The *image* of the *same* vertex from a *different* point of view gives a *different* junction type

(from *Winston, Intro to Artificial Intelligence*)

L Junctions have *only* 6 interpretations!



The world *constrains possibilities*

<i>Type of Junction</i>	<i>Physically Possible Interpretations</i>	<i>Combinatoric Possibilities</i>
<i>Arrow</i>	3	$4 \times 4 \times 4 = 64$
<i>L</i>	6	$4 \times 4 = 16$
<i>T</i>	4	$4 \times 4 \times 4 = 64$
<i>Y</i>	3	$4 \times 4 \times 4 = 64$



Idea 3 (*big idea*): Inference in CSPs

- **CSP solvers combine search *and* inference**
 - Search
 - assigning a value to a variable
 - *Constraint propagation (inference)*
 - Eliminates possible values for a variable if the value would violate *local consistency*
 - *Can do inference first, or intertwine it with search*
 - You'll investigate this in the Sudoku homework
- **Local consistency**
 - *Node consistency*: satisfies unary constraints
 - This is trivial!
 - *Arc consistency*: satisfies binary constraints
 - X_i is arc-consistent w.r.t. X_j if for every value v in D_i , there is some value w in D_j that satisfies the binary constraint on the arc between X_i and X_j .

An Example Constraint: The Edge Consistency Constraint


The Edge Consistency Constraint

Any *consistent assignment* of labels to the junctions in a picture must assign the *same* line label to any given line.




An Example of Edge Consistency

- Consider an arrow junction with an L junction to the right:



- A1 and either L1 or L6 are consistent*** since they both associate the same kind of arrow with the line.




- A1 and L2 are inconsistent***, since the arrows are pointed in the opposing directions,



- Similarly, A1 and L3 are inconsistent.***



Replacing Search: Constraint Propagation Invented...



Dave Waltz's insight for line labeling:


- Pairs of *adjacent* junctions (junctions connected by a line) *constrain* each other's interpretations!
- By *iterating* over the graph, the edge-consistency *constraints* can be *propagated* along the connected edges of the graph.
- **Search:** Use constraints to *add* labels to find *one* solution
- **Constraint Propagation:** Use constraints to *eliminate* labels to simultaneously find *all* solutions

The Waltz/Mackworth Constraint

Propagation Algorithm for line labeling

1. Assign every junction in the picture a set of *all* Huffman/Clowes junction labels for that junction type;
2. Repeat until there is no change in the set of labels associate with any junction:
 3. For each junction i in the picture:
 4. For each neighboring junction j in the picture:
 5. Remove any junction label from i for which there is no edge-consistent junction label on j .

Waltz/Mackworth: An example




Inefficiencies: Towards AC-3

1. At each iteration, we only need to examine those X_i *where at least one neighbor of X_i has lost a value in the previous iteration.*
2. If X_i loses a value only because of edge inconsistencies with X_j , we *don't need to check X_j on the next iteration.*
3. Removing a value on X_i can only make X_j edge-inconsistent is with respect to X_i itself. Thus, we only need to check that *the labels on the pair (j,i) are still consistent.*

These insights lead a much better algorithm...


Directed arcs, Not undirected edges

- Given a pair of nodes X_i and X_j in the constraint graph connected by a constraint *edge*, we represent this not by a single undirected edge, but a *pair of directed arcs*.
- For a connected pair of junctions X_i and X_j , there are *two* arcs that connect them: (i,j) and (j,i) .



Arc consistency: the general case

- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y in Y



Arc Consistency

An arc (i,j) is **arc consistent** if and only if every value v on X_i is consistent with some label on X_j .

- To make an arc (i,j) arc consistent,
for each value v on X_i ,
if there is no label on X_j consistent with v then
remove v from X_i

When to Iterate, When to Stop?

The crucial principle:

*If a value is removed from a node X_i ,
then the values on all of X_i 's neighbors must be
reexamined.*

Why? *Removing* a value from a node may result in
one of the neighbors becoming arc *inconsistent*,
so we need to check...

(but each neighbor X_j can only become inconsistent
with respect to the removed values on X_i)

AC-3

function **AC-3**(*csp*) return the CSP, possibly with reduced domains

inputs: *csp*, a binary csp with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs initially the arcs in *csp*

while *queue* is not empty do

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

 if **REMOVE-INCONSISTENT-VALUES**(X_i, X_j) then

 for each X_k in **NEIGHBORS**[X_i] – $\{X_j\}$ do

 add (X_k, X_i) to queue

function **REMOVE-INCONSISTENT-VALUES**(X_i, X_j) return *true* iff we remove a value

removed \leftarrow *false*

 for each *x* in **DOMAIN**[X_i] do

 if no value *y* in **DOMAIN**[X_j] allows (*x*,*y*) to satisfy the constraints between X_i and X_j

 then delete *x* from **DOMAIN**[X_i]; *removed* \leftarrow *true*

 return *removed*

AC-3: Worst Case Complexity Analysis

- All nodes can be connected to *every* other node,
 - so each of n nodes must be compared against $n-1$ other nodes,
 - so total # of arcs is $2*n*(n-1)$, i.e. $O(n^2)$
- If there are d values, checking arc (i,j) takes $O(d^2)$ time
- Each arc (i,j) can only be inserted into the queue d times
- Worst case complexity: $O(n^2d^3)$

For *planar* constraint graphs, the number of arcs can only be *linear in N* , so for our pictures, the time complexity is only $O(nd^3)$

- *The constraint graph for line drawings is isomorphic to the line drawing itself, so is a planar graph.*


Beyond binary constraints:

Path consistency

- Generalizes arc-consistency from individual binary constraints to multiple constraints
- A pair of variables X_i, X_j is path-consistent w.r.t. X_m if for every assignment $X_i=a, X_j=b$ consistent with the constraints on X_i, X_j there is an assignment to X_m that satisfied the constraints on X_i, X_m and X_j, X_m
- **Global constraints**
 - Can apply to any number of variables
 - E.g., in Sudoku, all numbers in a row must be different
 - E.g., in cryptarithmic, each letter must be a different digit
 - Example algorithm:
 - If any variable has a single possible value, delete that variable from the domains of all other constrained variables
 - If no values are left for any variable, you found a contradiction


Chronological backtracking

- **DFS does Chronological backtracking**
 - If a branch of a search fails, backtrack to the most recent variable assignment and try something different
 - But this variable may not be related to the failure
- **Example: Map coloring of Australia**
 - Variable order
 - Q, NSW, V, T, SA, WA, NT.
 - Current assignment:
 - Q=red, NSW=green, V=blue, T= red
 - SA cannot be assigned anything
 - But reassigning T does not help!



Backjumping: Improved backtracking

- **Find “the conflict set”**
 - Those variable assignments that are in conflict
 - Conflict set for SA: {Q=red, NSW=green, V=blue}
- **Jump back to reassign one of those conflicting variables**
- **Forward checking can build the conflict set**
 - When a value is deleted from a variable’s domain, add it to its conflict set
 - But backjumping finds the same conflicts that forward checking does
 - Fix using “conflict-directed backjumping”
 - Go back to predecessors of conflict set




Local search for CSPs

- Hill-climbing, simulated annealing typically work with "complete" states, i.e., all variables assigned
- To apply to CSPs:
 - allow states with unsatisfied constraints
 - operators **reassign** variable values
- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts** heuristic:
 - choose value that violates the fewest constraints
 - i.e., hill-climb with $h(n)$ = total number of violated constraints

Example: n-queens

- **States:** 4 queens in 4 columns ($4^4 = 256$ states)
- **Actions:** move queen in column
- **Goal test:** no attacks
- **Evaluation:** $h(n)$ = number of attacks



- **Given random initial state, local min-conflicts can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)**



Simple CSPs can be solved quickly

1. Completely independent subproblems

- e.g. Australia & Tasmania
- Easiest

2. Constraint graph is a tree

- Any two variables are connected by only a single path
- Permits solution in time linear in number of variables
- Do a topological sort and just march down the list



Simplifying hard CSPs: Cycle Cutsets

- **Constraint graph can be decomposed into a tree**
 - Collapse or remove nodes
 - *Cycle cutset* S of a graph G : any subset of vertices of G that, if removed, leaves G a tree
- **Cycle cutset algorithm**
 - Choose some cutset S
 - For each possible assignment to the variables in S that satisfies all constraints on S
 - Remove any values for the domains of the remaining variables that are not consistent with S
 - If the remaining CSP has a solution, then you have are done
 - For graph size n , domain size d
 - Time complexity for cycle cutset of size c :
 $O(d^c * d^2(n-c)) = O(d^{c+2}(n-c))$